

Supplementary Material:

Overcoming high nanopore basecaller error rates for DNA storage via basecaller-decoder integration and convolutional codes

Shubham Chandak, Joachim Neu, Kedar Tatwawadi, Jay Mardia, Billy Lau, Matthew Kubit, Reyna Hulett, Peter Griffin, Mary Wootters, Tsachy Weissman, Hanlee Ji
Stanford University

Contents

1 Availability of code and data	1
2 Convolutional codes	1
2.1 Code Parameters	1
2.2 Start/end code state and impact of runs	2
2.3 Puncturing for higher rates	3
2.4 Handling reverse complemented reads	4
2.5 Removal of primers for convolutional code decoding	4
2.6 List decoding and implementation details	4
3 Reed Solomon outer code	4
4 CRC	5
5 Simulations	5
6 Experimental parameters and results	5
6.1 Impact of list size	8
6.2 Impact of Reed Solomon outer code redundancy	9
6.3 Results for previous works	9
7 Experimental procedure	9

1 Availability of code and data

The code and instructions for installation and encoding/decoding are available at https://github.com/shubhamchandak94/nanopore_DNA_storage/. The specific scripts used for running various experiments are mentioned in the README and the text below as relevant. The data (encoded files, oligo files sent for synthesis, raw sequencing data, error statistics, decoded lists) is available at https://github.com/shubhamchandak94/nanopore_DNA_storage_data/.

2 Convolutional codes

2.1 Code Parameters

We use three rate $1/2$ convolutional codes with memory $m = 8, 11, 14$ for our experiments. The code parameters were derived from [1] and are given below. The current implementation also supports another

code with $m = 6$ which is obtained from [2]. Note that a rate 1/2 convolutional code can be specified using the two output functions, which represent the linear combination of the state and current input used to obtain the output. Following the convention in [1], the output linear combinations (i.e., generator polynomials) are written in octal.

- $m = 6$: $G = [171, 133]$
- $m = 8$: $G = [515, 677]$
- $m = 11$: $G = [5537, 6131]$
- $m = 14$: $G = [75063, 56711]$

To better relate the parameters to the encoding circuit for $m = 6$, note that G can be written in binary as $[1111001, 1011011]$ where the 1's show the position of the connections in the encoding circuit in Figure 1. Note that the two output streams are interleaved, i.e., if the first output stream (shown as the top output) is $(C_1(1), C_1(2), C_1(3), \dots)$ and the second output stream (shown as the bottom output) is $(C_2(1), C_2(2), C_2(3), \dots)$, then the overall output is $(C_1(1), C_2(1), C_1(2), C_2(2), C_1(3), C_2(3), \dots)$




Figure 1: Encoding circuit for $m = 6$ code with $G = [171, 133]$.

2.2 Start/end code state and impact of runs

The convolutional code encoding typically starts and ends with a fixed and known state (usually the all zeros state) [3]. This is achieved by padding the input with m known bits at the end, thus guaranteeing a specific end state. For this work, we found that starting/ending with the all zero state was suboptimal because of synchronization issues. For example if the input bit sequence is 000111...1001011, the sequence of states starting from the all zeros state (for $m = 6$) would be 000000, 000000, 000000, 000000, 100000, 110000, Due to the repeated sequence of states at the beginning, the decoder tends to incorrectly predict the number of 0s at the start. This is because the number of raw signal samples corresponding to each base is random, a fact that leads to difficulties in inferring the length of repeated state sequences. The result is that the decoder tends to produce a shifted message. To resolve this issue, we adopt start and end states that do not lead to repeated states. The start states used are listed below (the end states were chosen as the reverse of the start states for symmetry in case of reverse complemented reads).

- $m = 6$: 100101
- $m = 8$: 10010110
- $m = 11$: 10010110001
- $m = 14$: 10010110001101

We face similar issues when the input bit sequence had a run of 0s or 1s of length bigger than m , leading to repeated states. As m increases, such runs are rarer (since the compression and encryption compressed make the input appear random), leading to improved performance. Furthermore, the situation improves with increasing list size as the correct run length is likely to be present somewhere in the list.

For the decoding process, the state includes the convolutional code state, the flappie basecaller state and the position in the input message. The most likely state sequences that end in a valid state is obtained using Viterbi decoding. The valid end state must have the convolutional code state as the end state and the position as the last position. When the position is not included in the state, we found that the optimal state sequences produced significantly shorter message lengths, thus we decided to include the position in the state to enforce the correct message length in the decoded output.

2.3 Puncturing for higher rates

To increase the convolutional code rate beyond $1/2$, we use the technique of puncturing [2]. The idea is to transmit only a prespecified subsequence of the bits, usually dictated by a repeating pattern. For example, suppose that $C_1(i)$ and $C_2(i)$ represent the two output streams of the code. Then the output without puncturing is given by $(C_1(1), C_2(1), C_1(2), C_2(2), C_1(3), C_2(3) \dots)$. Now, if we use a puncturing pattern

$C_1 : 0 \ 1$

$C_2 : 1 \ 0$

where 1 represents a transmitted symbol and 0 represents a non-transmitted symbol, then the output would look like $(C_2(1), C_1(2), C_2(3), C_1(4) \dots)$ and the new code rate is 1, i.e., one output bit per input bit. The current implementation uses the following patterns (generally based on [2] with some modifications due to reasons described below).

- Rate = $1/2$
 $C_1 : 1$
 $C_2 : 1$
- Rate = $2/3$
 $C_1 : 1 \ 1 \ 0 \ 1$
 $C_2 : 1 \ 0 \ 1 \ 1$
- Rate = $3/4$
 $C_1 : 1 \ 0 \ 1$
 $C_2 : 1 \ 1 \ 0$
- Rate = $4/5$
 $C_1 : 1 \ 0 \ 0 \ 1 \ 1 \ 0 \ 0 \ 1$
 $C_2 : 1 \ 1 \ 1 \ 1 \ 0 \ 1 \ 1 \ 0$
- Rate = $5/6$
 $C_1 : 1 \ 0 \ 1 \ 1 \ 0$
 $C_2 : 1 \ 1 \ 0 \ 0 \ 1$
- Rate = $7/8$
 $C_1 : 1 \ 0 \ 0 \ 0 \ 1 \ 0 \ 1$
 $C_2 : 1 \ 1 \ 1 \ 1 \ 0 \ 1 \ 0$

The binary output is converted to a DNA sequence with a 2 bits per base mapping ($00 \rightarrow A$, $01 \rightarrow C$, $10 \rightarrow G$, $11 \rightarrow T$). During the decoding the state transitions of the basecaller (flappie) correspond to either a stay (no change in current base) or a shift by one base. When the convolutional code rate is $1/2$, each basecaller transition by one base corresponds to one state transition in the convolutional code (since each input bit produces one base). With puncturing, the situation is a bit more involved. To simplify things as much as possible, we only use puncturing patterns composed of the following building blocks:

$$\begin{array}{cccccccc} 1 & 0 & 1 & 1 & 0 & 0 & 0 & 1 & 1 \\ 1 & ' & 1 & 0 & ' & 0 & 1 & ' & 1 & 1 & ' & 0 & 0 \end{array}$$

In each of these cases, either one or two input bits produce two output bits (i.e., one base). When one input bit produces one base, an example state transition in the decoder would be from $(010011, A+, 42)$ to $(001001, C+, 43)$ where the convolutional code state shifts by one. In this case, there are three transitions out of the state (one stay, and two corresponding to input bit 0 or 1). When two input bits produce one base, an example state transition in the decoder would be from $(010011, A+, 42)$ to $(000100, C+, 43)$ where the convolutional code state shifts by two. In this case, there are five transitions out of the state (one stay, and four corresponding to a pair of input bits). Finally, note that in some cases the input needs to be padded by one bit to make sure that the puncturing pattern is composed of the given building blocks.

2.4 Handling reverse complemented reads

The sequencing process involves double stranded DNA where the sequence or its reverse complement can be sequenced. The reverse complement sequence consists of the original sequence in reverse, where each base is replaced by its complement ($A \leftrightarrow T$, $C \leftrightarrow G$). The decoding of reverse complemented reads is based on the fact that reversing a codeword for a convolutional code produces a codeword for a different convolutional code which has the same minimum distance properties but different generator polynomials. In fact, the generator polynomials and the puncturing patterns are effectively reversed for the new convolutional code. Thus, the handling of reverse complemented reads is quite straightforward as long as we can detect whether the read is reverse complemented. The detection is done based on the primers as discussed in the next section.

2.5 Removal of primers for convolutional code decoding

The raw current signal obtained from the nanopore corresponds to a DNA sequence consisting of the PCR primers and sequencing adapters in addition to the encoded sequence. However, the decoder only expects the raw signal for the encoded sequence. To identify the part of the raw signal corresponding to the encoded sequence, we first perform basecalling using flappie. Then we search for the start and end primer in the basecalled sequence based on minimum edit distance. Based on the position of the primers and a correspondence table (obtained from flappie) relating the base position to the raw signal position, we truncate the transition probability array that the decoder receives. To detect reverse complemented reads, we perform the minimum edit search for the primers in both orientations (forward and reverse complemented) and pick the one that achieves lower edit distance.

2.6 List decoding and implementation details

We use the parallel List Viterbi Algorithm from [4] with certain modifications. The list algorithm attempts to find the L most likely paths instead of just the most likely path. This is done by keeping track of the top L paths entering each state at each time step. We faced two major issues with this: (i) since multiple state sequences can correspond to the same message (due to stay transitions), finding the top L best paths through the states tends to produce a list containing just 1 or 2 distinct messages (rather than L), (ii) the memory consumption grows linearly with L and rapidly becomes intractable, chiefly due to the traceback array which needs to store L entries (previous state) for each state at each time step.

To resolve these issues, we utilize the fact that we do not need to know the most likely state sequence, instead we are interested in the messages themselves. Thus, instead of using a traceback array that needs to store the previous state corresponding to the best path at each state at *each* time step, we just store a list of L most likely messages along with their scores at each state at the *current* time step. Storing the messages at the current state also allows us to keep only the distinct messages at each state. The implementation uses heaps to perform a merge of the incoming lists at each state to find out the list of L most likely distinct messages at this state.

Finally, to further optimize the decoding, we do not allow states for which the position in the codeword deviates significantly from the expected position based on the location in the raw signal. Allowing a large enough deviation (20) still provides significant speedup while having essentially no impact on the optimality. The implementation is parallelized and can operate on the different states at a given timestep in parallel. If sufficient memory is available, further speedup can be gained by running the decoder on different reads in parallel.

3 Reed Solomon outer code

We use Reed Solomon (RS) code as the outer code to recover lost sequences and to correct any erroneous sequences which were not detected by the CRC. The RS code implementation is similar to that in [5] and is built on the Schifra library [6].

Figure 2 illustrates the Reed Solomon encoding procedure. The input data is segmented into segments of length $16n_{RS}$ where n_{RS} is an integer. The length of the segment is decided based on the desired length of the synthesized oligos after the inner coding. The RS coding with symbol size of 16 bits is applied independently




Figure 2: Reed Solomon encoding schematic.

to the n_{RS} chunks of the segments as shown in the figure, generating new parity segments. If the number of data segments is n_{data} and the RS redundancy is f (30% by default), then the number of parity segments $n_{parity} = f n_{data}$. Since we are working with symbol size of 16 bits, $n_{data} + n_{parity} < 65536$, which means that $n_{data} < n_{data}^{MAX} := 65536/(1 + f)$. For larger data sizes, we need to apply the procedure independently to groups of n_{data}^{MAX} segments.

Smaller symbol sizes lead to very few segments in each block, leading to high probability that some of blocks fail in the decoding (due to channel dispersion effects). Larger symbol sizes lead to higher computational complexity and also lack good existing implementations. The choice of 16 as the symbol size was made to tradeoff these factors. To understand why RS coding is applied to vertical blocks as shown in Figure 2, note that the erasures always occur at the segment-level. The vertical strategy allows erasure coding over a larger number of segments (up to 65,535 including parity segments) and similar strategies are often applied for bursty errors (interleaving).

During decoding, the set of segments decoded by the inner code are used as input to the RS code decoder. For each block, the decoding succeeds as long as $n_{erasure} + 2n_{error} < n_{parity}$ where $n_{erasure}$ is the number of erased (lost) segments and n_{error} is the number of erroneous segments (due to CRC detection failure). An increase in n_{error} leads to the RS code being able to handle fewer erasures which significantly increases the reading cost, especially in the presence of coverage bias [7]. Note that one error requires twice the amount of RS redundancy than one erasure and hence there is a need for error detection, as described next.

4 CRC

We use a 8-bit CRC for error detection after the convolutional code decoding and also to assist in the list decoding. The short CRC length was chosen to reduce the writing and reading cost. The `crc8` library in Python3 (<https://pypi.org/project/crc8/>) was used for the implementation. In cases where we get multiple reads with the same index that satisfy the CRC, we use the most commonly occurring sequence among those for the outer code decoding.

5 Simulations

Several of the parameters for the convolutional code were chosen based on simulations (code: `simulator.py`). We tried two simulators, `scrapie` [8] and `DeepSimulator` [9]. Overall, we found that `scrapie` provided results closer to what we observed in practice, even though it's dwell time is quite idealistic and has low variance. Therefore, we provide an option to use the dwell times from `scrapie` or `DeepSimulator` while using the signal levels (mean and variance) from `scrapie`.

6 Experimental parameters and results

The input file used for our experiments consisted of the following text files:

- gettysburg.txt: The Gettysburg Address by Abraham Lincoln
- mlkjr.txt: “I have a Dream” by Martin Luther King, Jr
- poems.txt: A collection of poems:
 - “The Road Not Taken” by Robert Frost
 - “Stopping by Woods on a Snowy Evening” by Robert Frost
 - “If” by Rudyard Kipling
 - “Fire and Ice” by Robert Frost
 - “The Tyger” by William Blake
 - “A Psalm of Life” by Henry Wadsworth Longfellow
 - “Sonnet 18: Shall I compare thee to a summer’s day?” by William Shakespeare
 - “I Wandered Lonely as a Cloud” by William Wordsworth
 - “To Autumn” by John Keats
- rickastley.txt: Lyrics of “Never Gonna Give You Up” by Rick Astley
- unhr.txt: The Universal Declaration of Human Rights

These are available at https://github.com/shubhamchandak94/nanopore_dna_storage_data/tree/master/encoded_file. The files were tarred and compressed with bzip2 and then encrypted to avoid homopolymers. The final file size was 11,280 bytes.

m	r	RS redundancy	#oligos data	#oligos RS	#oligos total	Payload bits per oligo	Convolutional code input length	Oligo length	Writing cost (bases/bit)	#reads for decoding	Reading cost (bases/bit)
8	1/2	30%	1128	338	1466	80	100	108	0.91	5500	6.58
8	3/4	30%	627	188	815	144	164	111	0.91	8000	10.20
8	5/6	30%	564	169	733	160	180	114	1.75	13500	16.91
11	1/2	30%	1128	338	1466	80	100	115	1.04	4500	5.54
11	3/4	30%	627	188	815	144	164	117	0.92	6000	7.78
11	5/6	30%	564	169	733	160	180	119	1.80	5500	7.01
14	1/2	30%	1128	338	1466	80	100	113	1.06	3500	4.42
14	3/4	30%	627	188	815	144	164	115	0.93	5000	6.59
14	5/6	30%	564	169	733	160	181	117	1.85	6500	8.43
11	3/4	20%	627	125	752	144	164	117	1.07	5500	7.13
11	3/4	40%	627	250	877	144	164	117	0.95	5000	6.48

Table 1: Parameters for the experiments. The input to the convolutional code contains the index (12 bits) and the CRC (8 bits) in addition to the payload bits. For $m = 14, r = 5/6$ a padding of 1 bit was applied due to the boundary condition for the puncturing pattern. The oligo length and the writing/reading costs exclude the primer with length 25 bases on either end. The number of reads for decoding shown here correspond to list size of 8 for $m = 8, 11$ and list size of 4 for $m = 14$.

The input file was encoded into oligos for several experiments, as shown in Table 1. We varied the convolutional code memory m and the rate r , and also tested a few values of the RS outer code redundancy for $m = 11, r = 3/4$. We also performed two more experiments for sanity check purposes, but they were not used for decoding and are not listed in the table. Additional encoding parameters such as the primers and the code used for encoding are available in `encode.experiments.py`. Note that the writing cost is calculated as $(\text{file size in bytes} \times 8) / (\# \text{oligos total} \times \text{Oligo length})$ while the reading cost is calculated as $(\text{file size in bytes} \times 8) / (\# \text{reads for decoding} \times \text{Oligo length})$.

After synthesis and sequencing, the reads were basecalled using the default Guppy basecaller (by ONT) and aligned to the original sequences sent for synthesis. The basecalled reads had close to 10.5% error rate with around 3-4% insertions, deletions and substitutions. Most of the errors are due to nanopore sequencing, since Customarray synthesis error rate is below 1% based on previous work in [7]. These reads were then separated into the respective experiments and the corresponding raw signal data was used to perform the decoding. Note that we used alignment to separate the various experiments for simplicity, but they can be separated either physically using PCR or computationally using the distinct primers. To enable comparison across

experiments, only the reads that align to the original sequences were considered for decoding. In practice, the unaligned/low-quality reads will be filtered out based on the primer sequence or using the CRC. Out of a total of 3.4 million sequenced reads, around 30% were unaligned and around 8% were chimeric (i.e., multiple oligos ligated together - see Figure 3 for the read length distribution). The script `util/align_compute_stats.sh` was used for these steps and the code in `util/read_length_distribution.cpp` was used for finding the read length distribution.




Figure 3: Read length distribution of the basecalled reads. There is a small peak at twice the expected length indicating chimeric reads due to ligation.

While the experiments were underway, a new high accuracy version of Guppy basecaller was released. We performed basecalling with this and the alignment was better, leading to only 15% unaligned reads with average basecalled error rate of 8.4%. The error rate is still quite high and does not affect our main conclusions regarding the efficacy of basecaller-decoder integration. Note that the flappie basecaller used for the decoding of the raw signal already uses the high accuracy model, so the main results concerning the reading cost for this approach are independent of this update.

m	r	RS redundancy	Writing cost (bases/bit)	Reading cost (bases/bit)						
				$L = 1$	$L = 2$	$L = 4$	$L = 8$	$L = 16$	$L = 32$	$L = 64$
8	1/2	30%	1.75	7.18	7.18	6.58	6.58	5.98	6.58	6.58
8	3/4	30%	1.04	12.74	11.47	11.47	10.20	10.83	12.11	14.02
8	5/6	30%	0.92	20.66	18.78	18.16	16.90	19.41	21.29	25.04
11	1/2	30%	1.80	6.15	6.15	5.54	5.54	-	-	-
11	3/4	30%	1.06	9.08	7.78	7.78	7.78	-	-	-
11	5/6	30%	0.93	8.92	8.28	7.01	7.01	-	-	-
14	1/2	30%	1.85	5.05	5.05	4.42	-	-	-	-
14	3/4	30%	1.07	7.25	7.25	6.59	-	-	-	-
14	5/6	30%	0.95	9.08	9.08	8.43	-	-	-	-
11	3/4	20%	0.98	8.43	7.78	7.78	7.13	-	-	-
11	3/4	40%	1.14	7.78	7.13	6.48	6.48	-	-	-

Table 2: Reading costs for different list sizes.

A random sample of raw signal files per experiment was then decoded with the convolutional code decoding (using `generate_decoded_lists.py`). The number of decoded reads was 20,000 for experiments with $m = 8$ and 10,000 for the rest (due to computational constraints). The list size was set to 64 for $m = 8$, 8 for $m = 11$ and 4 for $m = 14$. Note that the results for any list size lower than these can also be obtained by truncating the lists. This was followed by RS decoding of these lists for different list sizes using `decode_RS_from_decoded_lists.py`. For each list size, the minimum number of reads (in steps of 500) needed for successful decoding in 10 out of 10 subsampling trials was obtained and was used to compute the

reading costs shown in Table 2. Based on this, the default list size of 8 was chosen for $m = 8, 11$ and 4 for $m = 14$.

6.1 Impact of list size

From Table 2, we see that the reading cost typically reduces with the list size until a list size of 8, but then it starts increasing again. This is because the probability of having a random message in the list satisfying the CRC increases with the list size and hence the number of incorrectly decoded messages increases. This phenomenon is further explored in Table 3 which shows (for each experiment and list size), the percentage of reads that were (i) decoded correctly, (ii) had no CRC match in the list and, (iii) had a CRC match in the list that was incorrect. This was obtained using `compute_error_rate_from_decoded_lists.py`. Firstly, note that the percentage of correctly decoded reads increases with the list size, decreases with the code rate and increases with the convolutional code memory. This percentage is around 60%-80% for the $r = 1/2$ codes, around 30-40% for the $r = 3/4$ codes and around 20-30% for the $r = 5/6$ codes. As the list size increases to 64 for the $m = 8, r = 3/4$ code, the percentage of incorrect CRC matches grows as high as 9.33%, which means that around 18% of the RS redundancy is used up for the error correction leaving only $\sim 10\%$ for recovering missing oligos, which leads to higher reading costs in Table 2.

	L	1	2	4	8	16	32	64
$m = 8, r = 1/2$	% correct	59.65%	63.35%	66.33%	68.93%	70.99%	72.57%	73.98%
	% no CRC match	40.26%	36.49%	33.35%	30.47%	27.79%	25.40%	22.40%
	RS redundancy 30% % incorrect CRC match	0.06%	0.13%	0.29%	0.58%	1.20%	2.01%	3.60%
$m = 8, r = 3/4$	% correct	19.85%	22.52%	24.86%	26.90%	28.72%	30.23%	31.65%
	% no CRC match	79.75%	76.95%	74.34%	71.70%	68.77%	65.01%	60.05%
	RS redundancy 30% % incorrect CRC match	0.19%	0.33%	0.59%	1.19%	2.30%	4.55%	8.09%
$m = 8, r = 5/6$	% correct	15.40%	17.90%	19.98%	21.67%	23.53%	25.12%	26.54%
	% no CRC match	84.27%	81.61%	79.25%	76.87%	73.75%	69.85%	63.96%
	RS redundancy 30% % incorrect CRC match	0.17%	0.33%	0.61%	1.30%	2.55%	4.87%	9.33%
$m = 11, r = 1/2$	% correct	57.66%	59.91%	61.51%	62.80%	-	-	-
	% no CRC match	42.23%	39.94%	38.21%	36.60%	-	-	-
	RS redundancy 30% % incorrect CRC match	0.05%	0.09%	0.22%	0.54%	-	-	-
$m = 11, r = 3/4$	% correct	31.47%	34.87%	37.42%	39.62%	-	-	-
	% no CRC match	68.25%	64.66%	61.95%	59.30%	-	-	-
	RS redundancy 30% % incorrect CRC match	0.15%	0.34%	0.50%	0.95%	-	-	-
$m = 11, r = 5/6$	% correct	19.72%	22.39%	24.27%	25.91%	-	-	-
	% no CRC match	79.91%	77.05%	74.73%	72.47%	-	-	-
	RS redundancy 30% % incorrect CRC match	0.21%	0.40%	0.84%	1.46%	-	-	-
$m = 14, r = 1/2$	% correct	76.49%	78.66%	80.15%	-	-	-	-
	% no CRC match	23.39%	21.18%	19.60%	-	-	-	-
	RS redundancy 30% % incorrect CRC match	0.06%	0.10%	0.19%	-	-	-	-
$m = 14, r = 3/4$	% correct	39.23%	42.02%	44.17%	-	-	-	-
	% no CRC match	60.42%	57.55%	55.16%	-	-	-	-
	RS redundancy 30% % incorrect CRC match	0.15%	0.34%	0.50%	-	-	-	-
$m = 14, r = 5/6$	% correct	24.23%	26.40%	28.24%	-	-	-	-
	% no CRC match	75.43%	73.05%	70.88%	-	-	-	-
	RS redundancy 30% % incorrect CRC match	0.15%	0.36%	0.69%	-	-	-	-
$m = 11, r = 3/4$	% correct	35.94%	39.48%	42.03%	44.30%	-	-	-
	% no CRC match	63.90%	60.26%	57.38%	54.63%	-	-	-
	RS redundancy 20% % incorrect CRC match	0.1%	0.21%	0.54%	1.02%	-	-	-
$m = 11, r = 3/4$	% correct	29.82%	32.77%	35.14%	37.14%	-	-	-
	% no CRC match	70.03%	66.98%	64.40%	61.77%	-	-	-
	RS redundancy 40% % incorrect CRC match	0.09%	0.19%	0.40%	1.03%	-	-	-

Table 3: Percentage of reads correctly decoded, decoded with no CRC match found in list and incorrectly decoded with CRC match found. These add up to slightly below 100% since primer removal failed and convolutional code decoding was not performed for a very small percentage of reads (typically less than 0.2%).

6.2 Impact of Reed Solomon outer code redundancy

Based on the analysis in [7], we expect that increasing the outer code redundancy should increase the writing cost and decrease the reading cost. Looking at the results in Table 2 for $m = 11$, $r = 3/4$ with RS redundancy 20%, 30% and 40%, we see that the reading costs for 40% redundancy are the lowest which is as expected from the analysis. Interestingly, the 20% redundancy code has lower reading cost than the 30% redundancy code, but that might be due to randomness in the experiments and due to the higher number of correctly decoded reads for the 20% redundancy code (Table 3).

6.3 Results for previous works

We compared the results for our approach with two previous works [5, 10]. The works in [5] and [10] use the same encoding based on Reed Solomon outer coding and constrained coding to avoid homopolymers, but the decoder in [10] uses an improved consensus algorithm leading to significantly lower reading cost. The writing cost for these works was computed as the reciprocal of the metric “bits per base excluding primers” reported as 1.10 in [5]. Since both these works report coverage (36x and 22x, respectively) rather than reading cost, we use the formula reading cost = coverage \times writing cost.

Note that while we use the 2 bits per base mapping in this work, previous works [5, 10, 11] have explored other mappings with desirable properties such as lack of homopolymers (e.g., GGGG) which are quite important when working with the basecalled reads due to systematic and biased errors in basecalling. In fact, an experiment in [10] showed that the consensus-based decoding failed when the mapping was not used. Since we are not directly working with the basecalled sequence, we use the simple encoding that reduces the writing cost and makes the basecaller-decoder integration more straightforward. To avoid excessive homopolymers, we randomize the data using compression and encryption before the encoding process.

7 Experimental procedure

Oligonucleotides for all experiments were ordered as a single stranded DNA pool from GenScript. Pools were individually amplified using PCR primers specific to the experiment or subpool of interest using 1X Kapa HiFi PCR ReadyMix (Roche Biosystems), 500 nM PCR primers, and 1 microliter of the oligonucleotide pool under the following conditions: 98°C for 45 seconds; 12 cycles of 98°C for 15 seconds, 60°C for 30 seconds, and 72°C for 30 seconds; and a final extension step of 72°C for one minute. Each PCR reaction was purified with Ampure XP beads (Beckman Coulter) using standard protocols.

The thirteen pools were then quantified by fluorescence using the Qubit instrument (Thermo Fisher Scientific). They were then pooled and 300 femtomoles were used for downstream nanopore sequencing library preparation. Briefly, we performed End Repair and A-tailing of DNA products using the KAPA HyperPrep kit (Kapa Biosystems/Roche). The DNA was purified using Ampure XP beads (Beckman Coulter) under standard conditions. We then used the standard SQK-LSK109 library preparation from Oxford Nanopore Technologies to create DNA compatible for nanopore sequencing. Adapters containing motor protein was ligated to the DNA using 1X LNB buffer, AMX adapter, and KAPA HyperPrep ligase enzyme. The library was again purified using Ampure XP beads, but washed with SFB buffer and eluted in EB buffer (both Oxford Nanopore Technologies). The libraries were then loaded and sequenced on a MinION 9.4.1 flowcell for 48 hours using standard protocols.

References

- [1] G. Liva *et al.*, “Code Design for Short Blocks: A Survey,” *CoRR*, vol. abs/1610.00873, 2016.
- [2] “TM synchronization and channel coding – summary of concept and rationale,” Tech. Rep. 130.1-G-2, CCSDS SLS-C&S Working Group, November 2012.
- [3] A. Viterbi, “Convolutional codes and their performance in communication systems,” *IEEE Transactions on Communication Technology*, vol. 19, no. 5, pp. 751–772, 1971.

- [4] N. Seshadri and C.-E. Sundberg, “List Viterbi decoding algorithms with applications,” *IEEE transactions on communications*, vol. 42, no. 234, pp. 313–323, 1994.
- [5] L. Organick *et al.*, “Random access in large-scale DNA data storage,” *Nature biotechnology*, vol. 36, no. 3, p. 242, 2018.
- [6] “Schifra: C++ Reed Solomon Error Correcting Library.” <https://github.com/ArashPartow/schifra>. Last accessed: October 3, 2019.
- [7] S. Chandak *et al.*, “Improved read/write cost tradeoff in DNA-based data storage using LDPC codes,” *bioRxiv*, 2019.
- [8] “Scrappie: a technology demonstrator for the Oxford Nanopore Research Algorithms group.” <https://github.com/nanoporetech/scrappie>. Last accessed: October 3, 2019.
- [9] Y. Li *et al.*, “DeepSimulator: a deep simulator for Nanopore sequencing,” *Bioinformatics*, vol. 34, pp. 2899–2908, 04 2018.
- [10] R. Lopez *et al.*, “DNA assembly for nanopore data storage readout,” *Nature communications*, vol. 10, no. 1, p. 2933, 2019.
- [11] S. H. T. Yazdi *et al.*, “Portable and error-free DNA-based data storage,” *Scientific reports*, vol. 7, no. 1, p. 5011, 2017.